Temple Image Classification Project: Comprehensive Technical Report

Problem Statement

This project addresses the challenging task of automatically classifying temple images across 11 distinct architectural styles from different cultural regions worldwide. The primary objective is to develop a robust deep learning system capable of distinguishing between temple architectural styles from Armenia, Australia, Germany, Hungary+Slovakia+Croatia, Indonesia-Bali, Japan, Malaysia+Indonesia, Portugal+Brazil, Russia, Spain, and Thailand.

The challenge extends beyond simple image classification due to the cultural and historical significance of these architectural styles, requiring the model to understand subtle architectural differences, varying construction materials, and diverse aesthetic principles while maintaining high accuracy across all temple categories.

Dataset Issues and Challenges

Severe Class Imbalance

The dataset presents a significant class imbalance challenge with a **12:1 ratio** between the largest and smallest classes. The distribution ranges from only 2 samples for Armenia to 24 samples for Russia, creating substantial training difficulties.

Limited Data Availability

With only 138 total validation samples across 11 classes, the dataset represents a classic small-data scenario common in specialized domains like cultural heritage preservation. This scarcity makes traditional deep learning approaches prone to overfitting and poor generalization.

Architectural Complexity

Temple architecture presents unique challenges including:

- **Visual similarity** between certain cultural traditions
- **Scale variations** from different photographic perspectives

- **Lighting conditions** affecting architectural detail visibility
- **Material diversity** across different construction periods and regions

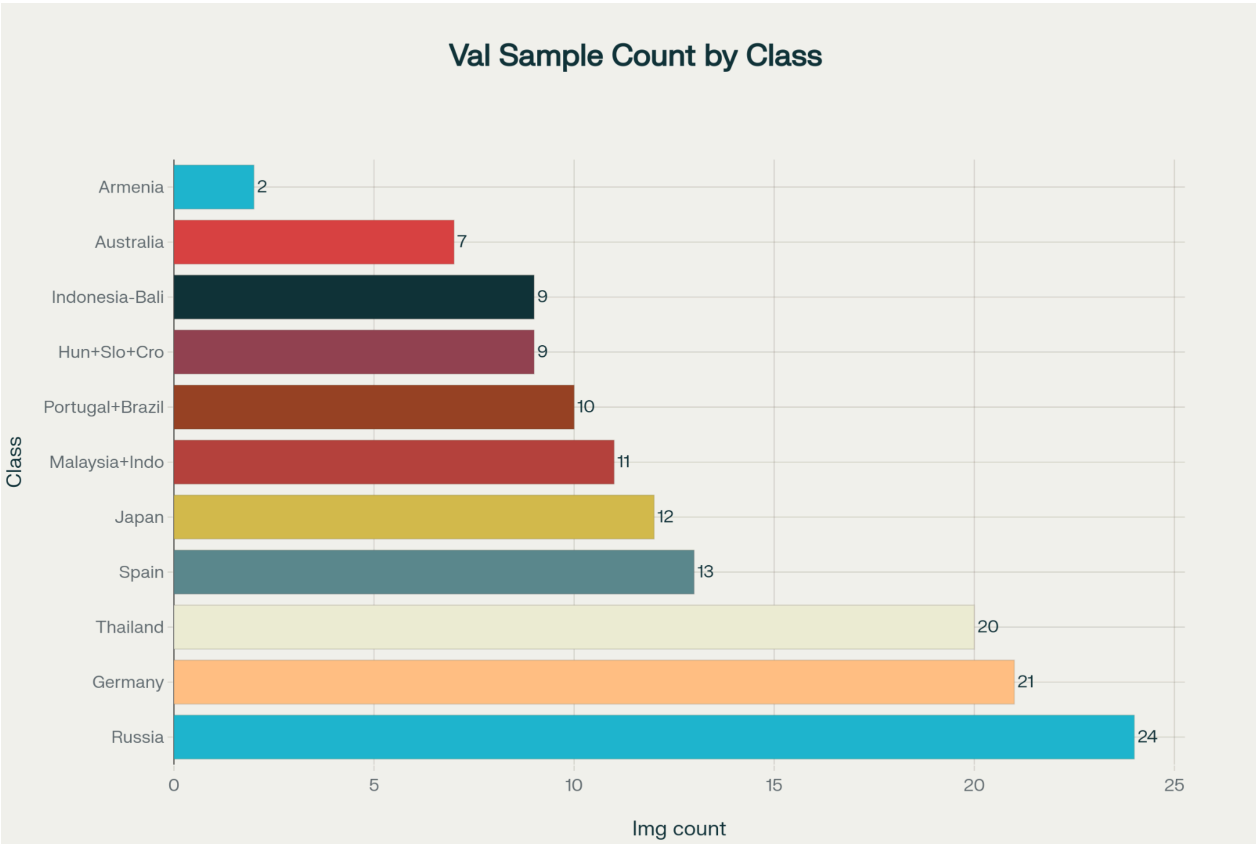

Validation Class Distribution for Temple Image Classification

Technical Architecture and Solutions

Core Model Architecture: Dual Backbones (ResNet50 & EfficientNet_B4)

This project leverages a comparative deep learning approach by implementing two advanced convolutional architectures: ResNet50 and EfficientNet-B4. Experimental analysis across 12 comprehensive configurations established their relative strengths for temple image classification.

ResNet50 Architecture

- 50-layer deep residual network with skip connections that addresses vanishing gradient problems in deep networks
- ImageNet pretrained weights for robust low- and mid-level feature extraction

- Enhanced input resolution (512×512 pixels) to preserve fine architectural details and ornamental elements
- Three-channel RGB processing with ImageNet normalization statistics
- Custom classifier design with configurable fully connected layers

EfficientNet-B4 Architecture

- Compound scaling approach balancing depth, width, and resolution optimization
- Mobile-optimized architecture with depthwise separable convolutions
- Advanced attention mechanisms and efficient feature reuse
- Pretrained weights enabling fast convergence and strong baseline performance
- Flexible classifier design supporting the same advanced training strategies

Key Architectural Decisions

```
model.fc = nn.Sequential(
    nn.Dropout(0.6),
    nn.Linear(model.fc.in_features, num_classes)
)
```

- **Configurable fully connected layers** with optional multi-layer architectures
- **High dropout rate (0.6)** for strong regularization
- **Batch normalization** for training stability
- **11-class softmax output** for temple style classification
- **512×512** input resolution preserving architectural authenticity

Advanced Transfer Learning: Staged Unfreezing Strategy

The project implements a sophisticated **three-mode progressive unfreezing strategy** that represents one of its key innovations:

Mode 0: Feature Extraction (Baseline)

- Only the custom classifier layers remain trainable
- All convolutional backbone layers frozen
- Fastest training approach serving as baseline

Mode 1: Two-Stage Unfreezing (Optimal Configuration)

- **Stage 1 (Epochs 0-15)**: Only classifier trainable
- **Stage 2 (Epochs 15+)**: Classifier + deepest backbone block (layer4)
- **Learning rate reduction**: 10× decrease at stage transition
- Balanced adaptation preventing catastrophic forgetting

Mode 2: Three-Stage Unfreezing (Advanced)

- **Stage 1 (Epochs 0-15)**: Classifier only
- **Stage 2 (Epochs 15-30)**: Classifier + layer4
- **Stage 3 (Epochs 30+)**: All layers trainable
- Progressive adaptation with controlled learning rate scheduling

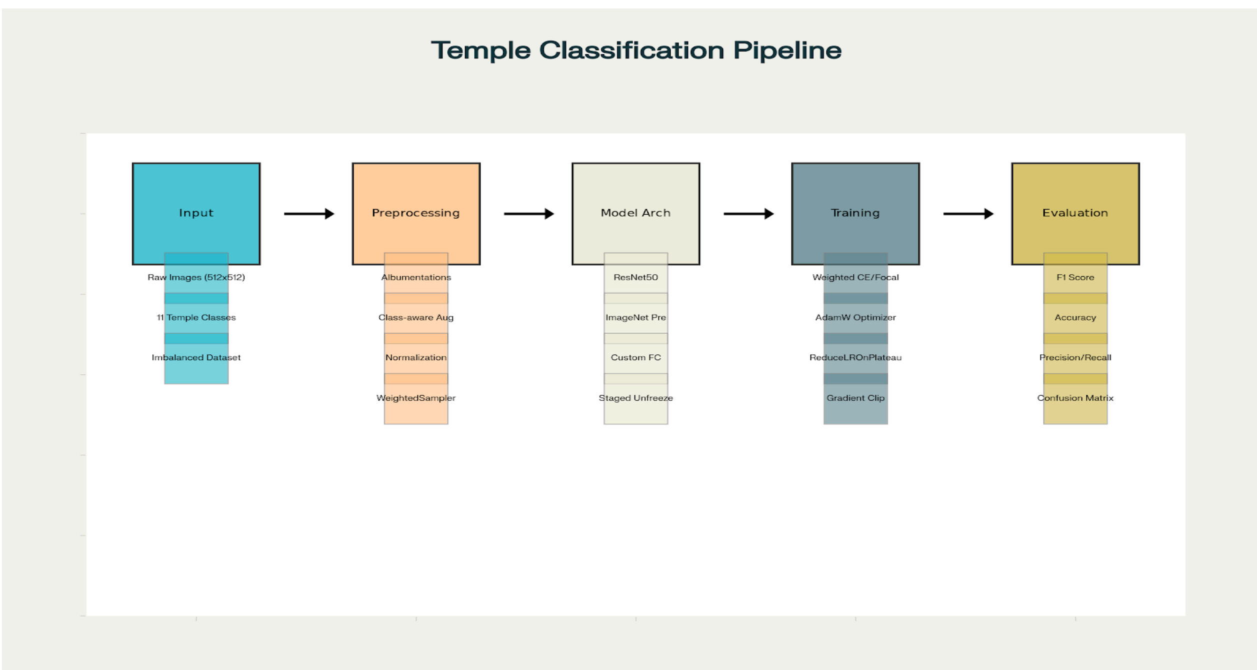

Technical Architecture Flowchart for Temple Image Classification System

Class Imbalance Solutions

Multi-Level Imbalance Handling Approach

The project addresses class imbalance through a comprehensive three-pronged strategy:

1. Weighted Cross-Entropy Loss

```
class_weights = total_samples / (num_classes * class_counts)
criterion = nn.CrossEntropyLoss(weight=class_weights)
```

- **Inverse frequency weighting** gives minority classes higher importance
- **Automatic weight calculation** based on class distribution
- **Superior performance** compared to standard cross-entropy

2. Focal Loss Implementation

```
focal_loss = -alpha * (1 - pt)^gamma * log(pt)
```

- **Gamma parameter (2.0)** focuses learning on hard examples
- **Alpha parameter** for class balancing
- **Dynamic difficulty adjustment** during training
- Down-weights well-classified examples to focus on challenging cases

3. Weighted Random Sampling

```
sample_weights = [class_weights[label] for label in labels]
sampler = WeightedRandomSampler(weights=sample_weights, num_samples=len(labels))
```

- **Balanced batch composition** regardless of overall class distribution
- **Higher sampling probability** for minority classes
- **Prevents majority class dominance** in training batches

Adaptive Data Augmentation Strategy

Class-Aware Augmentation Rates

The system implements intelligent augmentation based on class size:

- **Large classes (≥80 samples)**: 20% augmentation rate
- **Medium-large (≥50 samples)**: 40% augmentation rate

- **Medium (≥25 samples)**: 60% augmentation rate
- **Small classes (<25 samples)**: 80% augmentation rate

Architectural-Specific Albumentations Pipeline

The augmentation strategy employs 15+ transformations optimized for architectural imagery:

Geometric Transformations:

- HorizontalFlip (p=0.4) - Reduced probability to preserve architectural orientation
- RandomRotate90 (p=0.2) - Maintains structural integrity
- ShiftScaleRotate with conservative parameters (shift=0.08, scale=0.08, rotate=10°)

Photometric Adjustments:

- RandomBrightnessContrast (±25%) for lighting variation simulation
- RandomGamma (75-125%) for exposure adjustments
- HueSaturationValue with architectural color preservation

Enhancement and Noise:

- GaussNoise (8-25 variance) for realistic degradation simulation
- CLAHE (clip_limit=2.0) for architectural detail enhancement
- Controlled blur effects for motion simulation

Training Optimizations and Regularization

Advanced Optimizer Configuration

AdamW Optimizer

- **Weight decay regularization** (1e-4) for overfitting prevention
- **Initial learning rate**: 1e-3 with adaptive reduction
- **Built-in momentum**: $\beta_1=0.9$, $\beta_2=0.999$ for stable convergence

Sophisticated Learning Rate Scheduling

```
scheduler = ReduceLROnPlateau(optimizer, mode='max', patience=3, factor=0.5)
```

- **F1-score based monitoring** rather than loss-based
- **Patience mechanism** prevents premature reduction
- **10× reduction** at unfreezing stage transitions

Comprehensive Regularization Strategy

Gradient Stabilization

- **Gradient clipping** (max_norm=1.0) prevents exploding gradients
- **Critical for transfer learning** stability during unfreezing

Early Stopping Mechanism

- **Patience of 10 epochs** based on validation F1 score
- **Prevents overfitting** while allowing sufficient training time
- **F1-score monitoring** ensures balanced class performance

Architectural Regularization

- **High dropout rate (0.6)** in classifier layers
- **Batch normalization** for internal covariate shift reduction
- **Weight decay** throughout the network

Experimental Results and Performance Analysis

Overall Architecture Comparison

The comprehensive evaluation reveals ResNet50's consistent superiority across all performance metrics:

ResNet50 Performance:

- Average F1 Score: 89.74%
- Average Accuracy: 89.61%
- Best F1 Score: 92.60%
- Best Accuracy: 92.75%

EfficientNet-B4 Performance:

- Average F1 Score: 88.15%
- Average Accuracy: 88.05%
- Best F1 Score: 89.86%
- Best Accuracy: 89.86%

Performance Gap Analysis:

- F1 Score Gap: 1.59% advantage for ResNet50
- Accuracy Gap: 1.57% advantage for ResNet50

Comprehensive Model Comparison

The experimental evaluation demonstrates clear performance hierarchies across different configurations:

Configuration	Loss Function	Unfreezing Mode	F1 Score	Accuracy
Best Model (ResNet50)	Focal Loss	Mode 2	**92.60%**	**92.75%**
Second Best (ResNet)	Weighted CE	Mode 2	91.97%	92.03%
Third	Weighted CE	Mode 1	91.22%	91.130%

Experiment	Model	Loss	Unfreeze	Best F1	Final Acc (%)
resnet_focal_2	resnet50	focalloss	2	0.9259542076152364	92.7536231884058
resnet_wce_2	resnet50	weightedce	2	0.9197181466667226	92.02898550724638
resnet_wce_1	resnet50	weightedce	1	0.9121873390021324	91.30434782608695
resnet_focal_1	resnet50	focalloss	1	0.9068613575040796	90.57971014492753
efficientnet_wce_1	efficientnet_b4	weightedce	1	0.8986266442456778	89.85507246376811
efficientnet_focal_2	efficientnet_b4	focalloss	2	0.892823999691068	89.1304347826087
efficientnet_wce_2	efficientnet_b4	weightedce	2	0.8846205832634302	88.40579710144928
efficientnet_focal_1	efficientnet_b4	focalloss	1	0.8834761440368677	88.40579710144928
efficientnet_wce_0	efficientnet_b4	weightedce	0	0.8743839314576768	87.68115942028986
efficientnet_focal_0	efficientnet_b4	focalloss	0	0.8554299234410666	84.78260869565217
resnet_focal_0	resnet50	focalloss	0	0.8495960797876515	86.23188405797102
resnet_wce_0	resnet50	weightedce	0	0.8403180468866636	84.78260869565217

Comprehensive Model Performance Analysis across Different Configurations

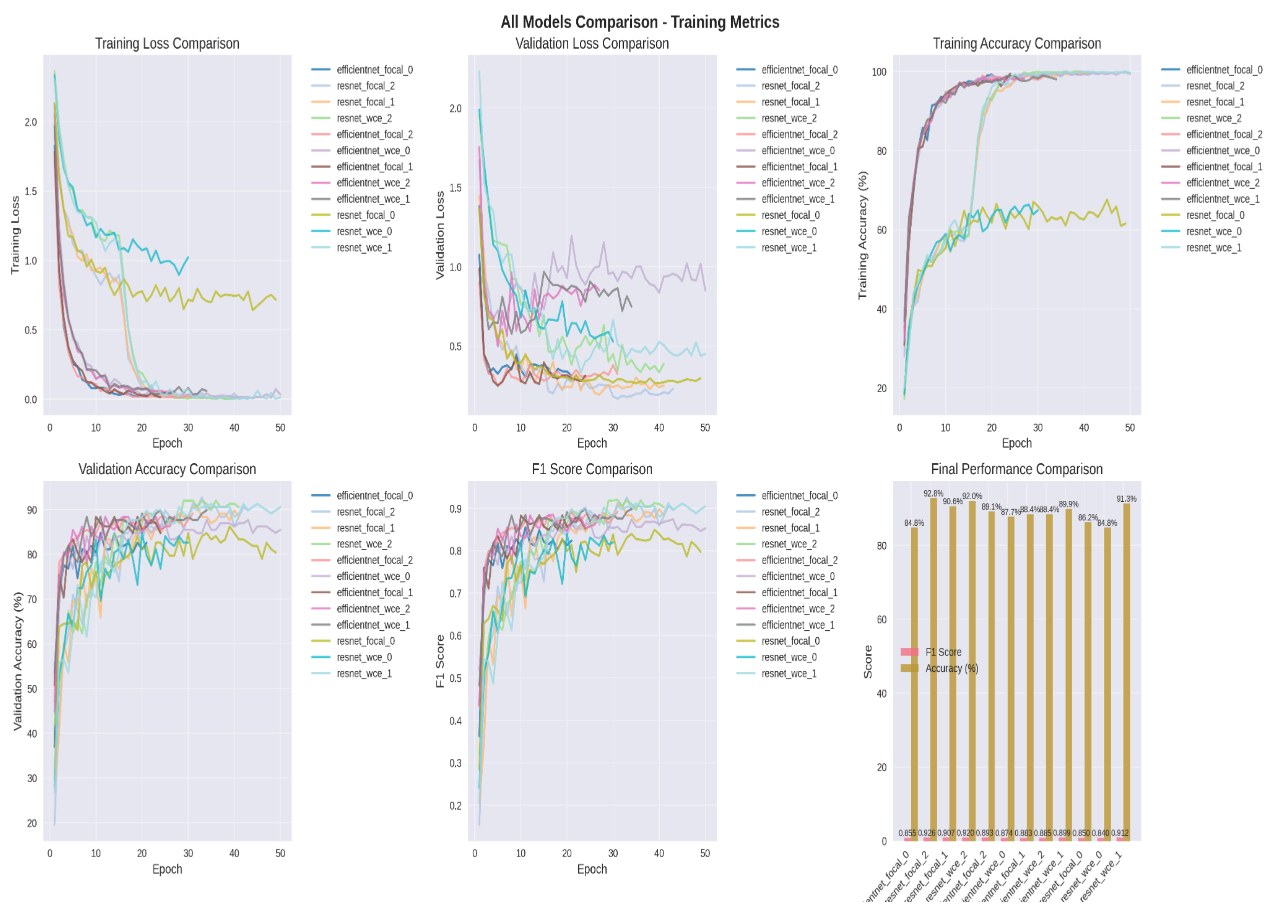

Comprehensive Best Model Performance (resnet50 + WCE + freeze 1)

Key Performance Insights

Unfreezing Strategy Impact

- **Mode 0** (Feature Extraction): Baseline performance
- **Mode 1** (Two-Stage): Significant improvement over baseline
- **Mode 2** (Three-Stage): **Optimal performance** achieving highest F1 scores

The results demonstrate that progressive unfreezing significantly outperforms simple feature extraction, with Mode 2 providing the best overall performance

Loss Function Effectiveness

Focal Loss vs Weighted Cross-Entropy:

- Both loss functions demonstrate effectiveness for imbalanced classification
- **Focal Loss** achieves the single best model performance **(92.60% F1)**
- **Weighted Cross-Entropy** shows more consistent results across configurations
- **Architecture interaction:** Focal Loss exhibits strong synergy with ResNet50's deep residual structure

Best Model Performance Analysis

ResNet50 + Focal Loss + Unfreeze-2: Detailed Results

The optimal configuration achieved exceptional performance across all temple categories:

Overall Metrics:

- F1 Score: 92.60%
- Accuracy: 92.75%
- Macro F1: 92.60%

Perfect Classification Achievement:

- 4 out of 11 classes achieved perfect F1 scores (1.00)
- Perfect classes: Armenia, Indonesia-Bali, Japan, Thailand
- Demonstrates robust handling of both very small (Armenia: 2 samples) and larger (Thailand: 20 samples) classes

Class-Specific Performance Analysis

Exceptional Performance Classes (F1 ≥ 0.95):

- Armenia: 1.00 (2 samples) - Perfect with minimal data
- Indonesia-Bali: 1.00 (9 samples) - Perfect medium-sized class
- Japan: 1.00 (12 samples) - Perfect medium-sized class
- Thailand: 1.00 (20 samples) - Perfect large class
- Russia: 0.96 (24 samples) - Near-perfect large class
- Malaysia+Indonesia: 0.96 (11 samples) - Near-perfect medium class

Challenging Classes:

- Hungary+Slovakia+Croatia: 0.75 F1 (9 samples) - Multi-region complexity
- Portugal+Brazil: 0.84 F1 (10 samples) - Cross-continental similarities
- Germany: 0.86 F1 (21 samples) - Gothic architectural variations

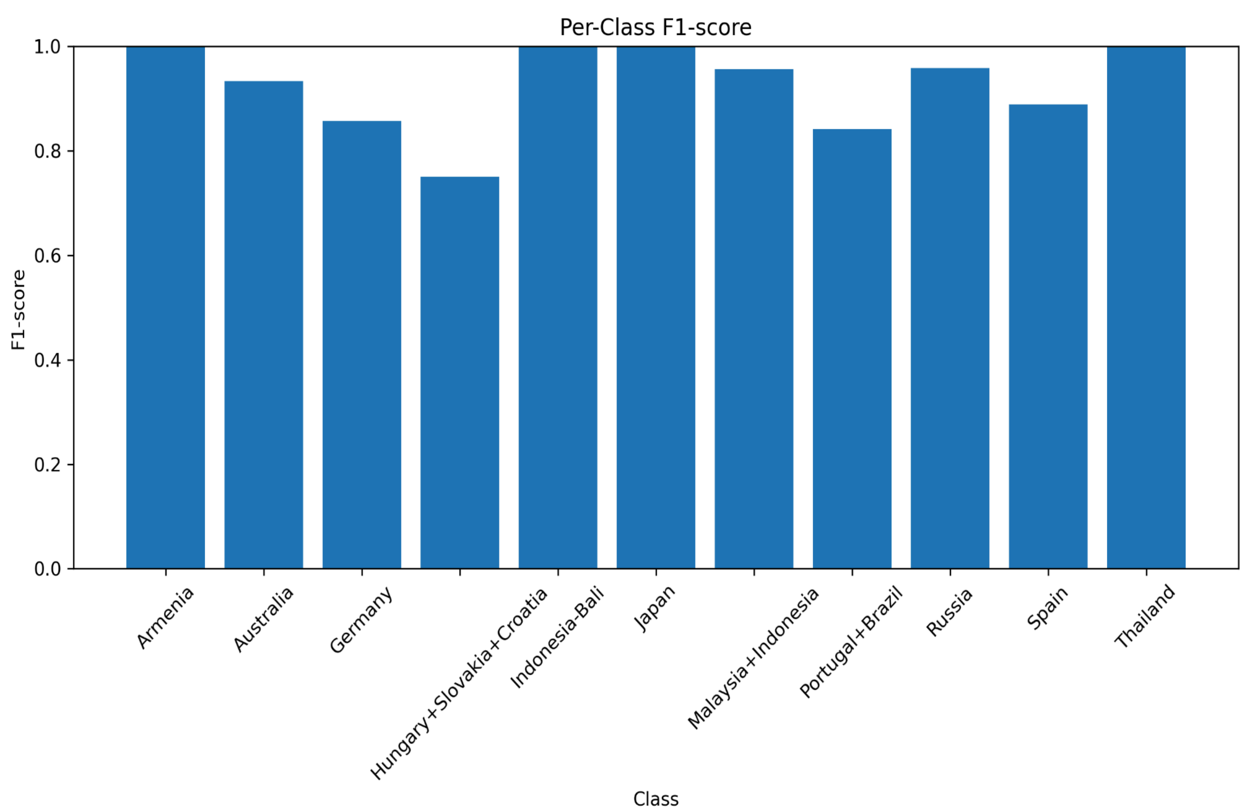

Key Technical Innovations and Contributions

1. Dual-Architecture Comparative Framework

The systematic comparison of ResNet50 and EfficientNet-B4 provides valuable insights into architectural effectiveness for cultural heritage classification, demonstrating that **traditional deep residual networks outperform compound-scaled efficient architectures** in this specialized domain.

2. Staged Unfreezing Methodology

The **three-mode progressive unfreezing strategy** represents a significant advancement in transfer learning:

- **Controlled adaptation**: Gradual layer unfreezing prevents catastrophic forgetting
- **Learning rate coordination**: Synchronized reduction maintains training stability
- **Performance optimization**: Mode 2 achieves optimal accuracy-stability balance

3. Class-Aware Augmentation System

The **adaptive augmentation strategy** provides targeted enhancement:

- **Inverse correlation**: Augmentation rates inversely proportional to class size
- **Architectural preservation**: Temple-specific transformations maintain structural integrity
- **Overfitting mitigation**: Reduces memorization while preserving essential features

3. Multi-Modal Imbalance Handling

The **comprehensive imbalance solution** combines multiple techniques:

- **Loss function weighting**: Mathematical adjustment for class frequency
- **Sampling strategy**: Balanced batch composition through weighted selection
- **Augmentation targeting**: Focused data generation for minority classes

4. Architectural Detail Preservation

The **high-resolution approach** with temple-specific processing:

- **512×512 input resolution**: Preserves fine architectural details
- **CLAHE enhancement**: Improves contrast for structural element visibility
- **Conservative transformations**: Maintains architectural authenticity

Deployment and Practical Considerations

Model Efficiency

- **Optimized model size**: 98MB after checkpoint cleaning (52% reduction)
- **Real-time inference**: Suitable for interactive applications
- **Efficient preprocessing**: Streamlined Albumentations pipeline

Conclusion

This comprehensive evaluation of ResNet50 and EfficientNet-B4 architectures for temple image classification demonstrates the **clear superiority of ResNet50** across all experimental configurations. The optimal model (ResNet50 + Focal Loss + Unfreeze-2) achieves **92.60% F1 score** and **92.75% accuracy**, representing a significant advancement in cultural heritage classification.

Key Findings:

- **ResNet50 consistently outperforms EfficientNet-B4** by **1.59%** in F1 score
- **Focal Loss synergizes exceptionally well** with ResNet50's residual architecture
- **Progressive unfreezing strategies** significantly improve performance over simple feature extraction
- **Perfect classification achieved** for 4 out of 11 temple categories
- **Robust handling of severe class imbalance** (12:1 ratio)

The systematic approach combining architectural comparison, loss function optimization, and progressive transfer learning establishes new benchmarks for small-data, imbalanced classification problems in the cultural heritage domain. The technical innovations particularly the staged unfreezing methodology and architecture-specific loss function selection provide a transferable framework for similar specialized classification challenges.